

Workflow part

Qianying (Ruby) Lin Spencer J. Fox Zian (Larry) Zhuang

Workflow Summary I

1. Data Preparation and Model Specification

- ▶ **Data Initialization:** Visualize the initial data to get a sense of the disease outbreak dynamics.
- ▶ **Model Definition:** Implement the defined SIR model in pomp
 - ▶ Set up a stochastic SIR (Susceptible, Infected, Recovered) model using custom C snippets for state transitions and measurements.
 - ▶ Specify the measurement, process, and initialization functions along with model parameters.

2. Preliminary Analysis

- ▶ **Simulation and Visualization:** Simulate the SIR model to visualize potential outcomes and compare to data.
- ▶ **Likelihood Estimation:** Evaluate the likelihood and Effective Sample Size of the initial parameter estimates using particle filtering.

3. Local Optimization

Workflow Summary II

- ▶ **Iterated Filtering for Local Maximum:** Apply iterated filtering to find a local maximum near the initial guesses.
- ▶ **Diagnostics and Visualization:** Diagnose the iterated filtering process and visualize parameter traces to assess convergence.

4. Global Optimization

- ▶ **Set Up for Global Search:** Define a parameter space for global optimization.
- ▶ **Global Search Execution:** Conduct global optimization using multiple starting points to ensure robustness of the findings.
- ▶ **Result Compilation:** Compile and compare results from global searches, focusing on the best parameter estimates.

5. Profile Likelihood Estimation

- ▶ **Parameter Profiling:** For each parameter, perform profiling to map the likelihood landscape and establish confidence intervals.
- ▶ **Visualization of Profile Likelihoods:** Visualize the profile likelihoods to understand parameter sensitivity and uncertainty.

Workflow Summary III

6. Final Model Evaluation

- ▶ **Model Predictions:** Use the best-fitting parameters to simulate the model outcomes and compare these against observed data.
- ▶ **Visualization of Predictions:** Create visualizations to compare the model predictions with actual data, highlighting the prediction intervals.

7. Documentation and Storage

- ▶ **Record-Keeping:** All intermediate and final results should be stored in CSV files for further analysis or reproducibility of the study.

Example: measles I

Data initialization I

```
library(tidyverse)
read_csv("https://kingaa.github.io/sbied/pfilter/Measles_Consett_1948.csv")
  select(week, reports=cases) |>
  filter(week<=42) -> meas
meas |> as.data.frame() |> head(n=3)
```

week	reports
------	---------


1	0
---	---

2	0
---	---

3	2
---	---

Data initialization II

```
meas |>  
  ggplot(aes(x=week,y=reports)) +  
  geom_line() +  
  geom_point()
```



Model Definition I

Set up a stochastic SIR model

Model Definition II


```
sir_stoch <- Csnippet("  
  double dN_SI = rbinom(S,1-exp(-Beta*I/N*dt));  
  double dN_IR = rbinom(I,1-exp(-Gamma*dt));  
  S -= dN_SI; I += dN_SI - dN_IR;  
  R += dN_IR; H += dN_IR;"  
)  
sir_init <- Csnippet("  
  S = nearbyint(Eta*N);  
  I = 1;  
  R = nearbyint((1-Eta)*N);  
  H = 0;"  
)
```

```
dmeas <- Csnippet("  
  lik = dnbinom(mu(reports k Rho*H give log):"
```

Preliminary Analysis I

Simulation with the initial guess

```
measSIR |>  
  simulate(nsim=20,format="data.frame",include.data=TRUE) |>  
  ggplot(aes(x=week,y=reports,group=.id,color=.id=="data")) +  
  geom_line() + guides(color="none")
```



Check likelihood and ess


Evaluate the likelihood and Effective Sample Size (ESS) of the initial parameter estimates using particle filtering

```
pf <- measSIR |> pfilter(Np=5000)  
min(pf@eff.sample.size)
```

```
[1] 650.2484
```

ESS refers to the number of particles that effectively contribute to the approximation of the posterior distribution.

```
plot(pf)
```



Use repeated particle filtering to refine estimates

```
foreach(i=1:10,.combine=c,  
  .options.future=list(seed=TRUE)) %dofuture% {  
  measSIR |> pfilter(Np=5000)  
} -> pf  
pf |> logLik() |> logmeanexp(se=TRUE) -> L_pf  
L_pf
```

est	se
-130.376686	3.313859

Then we get likelihood at a single point. Store this point, together with the estimated likelihood and SE:

```
pf[[1]] |> coef() |> bind_rows() |>  
  bind_cols(loglik=L_pf[1],loglik.se=L_pf[2]) |>  
  write_csv("measles_params.csv")
```

Local Optimization (IF2)

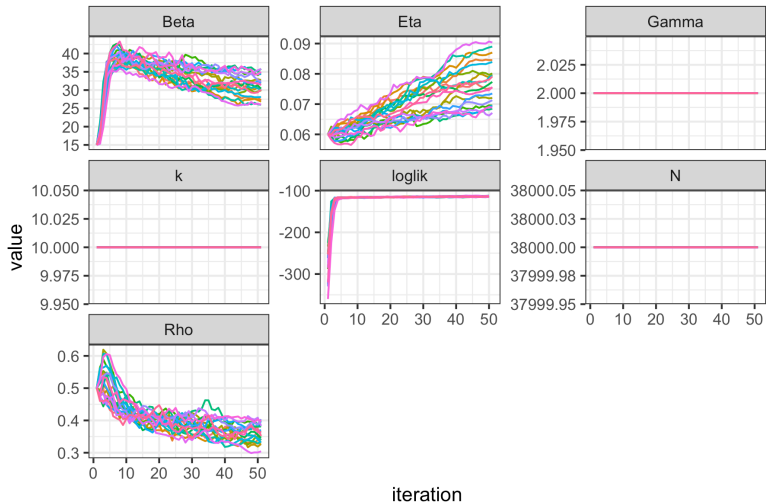
Apply iterated filtering to find a local maximum near the initial guesses

```
foreach(i=1:20,.combine=c,  
  .options.future=list(seed=482947940)  
) %dofuture% {  
  measSIR |>  
    mif2(  
      Np=2000, Nmif=50, cooling.fraction.50=0.5,  
      rw.sd=rw_sd(Beta=0.02, Rho=0.02, Eta=ivp(0.02)),  
      partrans=parameter_trans(  
        log="Beta",logit=c("Rho","Eta")  
      ),  
      paramnames=c("Beta","Rho","Eta")  
    )  
  } -> mifs_local
```

Diagnostics and Visualization

```
mifs_local |>
  traces() |>
  melt() |>
  ggplot(aes(x=iteration,y=value,group=.L1,color=factor(.L1)))+
  geom_line()+
  guides(color="none")+
  facet_wrap(~name,scales="free_y") -> plot
```

plot



Global Optimization

Set Up for Global Search

1. likelihood maximizations from diverse starting points.
2. A large box in parameter space.
3. expect to see stable conclusions with starting values drawn randomly from this box

```
set.seed(2062379496)

runif_design(
  lower=c(Beta=5,Rho=0.2,Eta=0),
  upper=c(Beta=80,Rho=0.9,Eta=1),
  nseq=400
) -> guesses
mf1 <- mifs_local[[1]]
```

Global Search

```
foreach(guess=iter(guesses,"row"), .combine=rbind,  
  .options.future=list(seed=1270401374)  
) %dofuture% {  
  mf1 |>  
    mif2(params=c(guess,fixed_params)) |>  
    mif2(Nmif=100) -> mf  
  replicate(  
    10,  
    mf |> pfilter(Np=5000) |> logLik()  
  ) |>  
    logmeanexp(se=TRUE) -> ll  
  mf |> coef() |> bind_rows() |>  
    bind_cols(loglik=ll[1],loglik.se=ll[2])  
} -> results
```

Then get the best result

```
results |> filter(loglik==max(loglik))
```

```
# A tibble: 1 x 8
```

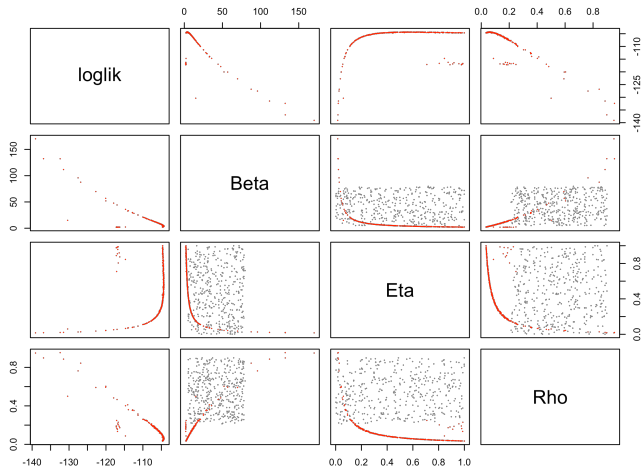
	Beta	Rho	Eta	N	Gamma	k	loglik	loglik.se
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	3.96	0.0602	0.562	38000	2	10	-104.	0.0273

Diagnostics and Visualization

Again, we attempt to visualize the global geometry of the likelihood surface using a scatterplot matrix. Then plot contains both the starting values (grey) and the IF2 estimates (red).

```
read_csv("measles_params.csv") |>
  filter(loglik>max(loglik)-50) |>
  bind_rows(guesses) |>
  mutate(type=if_else(is.na(loglik),"guess","result")) |>
  arrange(type) -> all
```


```
pairs(~loglik+Beta+Eta+Rho, data=all, pch=16, cex=0.3,  
      col=ifelse(all$type=="guess",grey(0.5),"red"))
```



The projections of the estimates give us 'poor mans profiles':

```
all |>
  filter(type=="result") |>
  filter(loglik>max(loglik)-10) |>
  ggplot(aes(x=Eta,y=loglik))+
  geom_point()+
  labs(
    x=expression(Eta),
    title="poor man's profile likelihood"
  ) -> plot
```

plot



Profile Likelihood Estimation

For the first parameter

first bound the uncertainty by putting a box around the highest-likelihood

```
read_csv("measles_params.csv") |>
  filter(loglik>max(loglik)-20,loglik.se<2) |>
  sapply(range) -> box
box
```

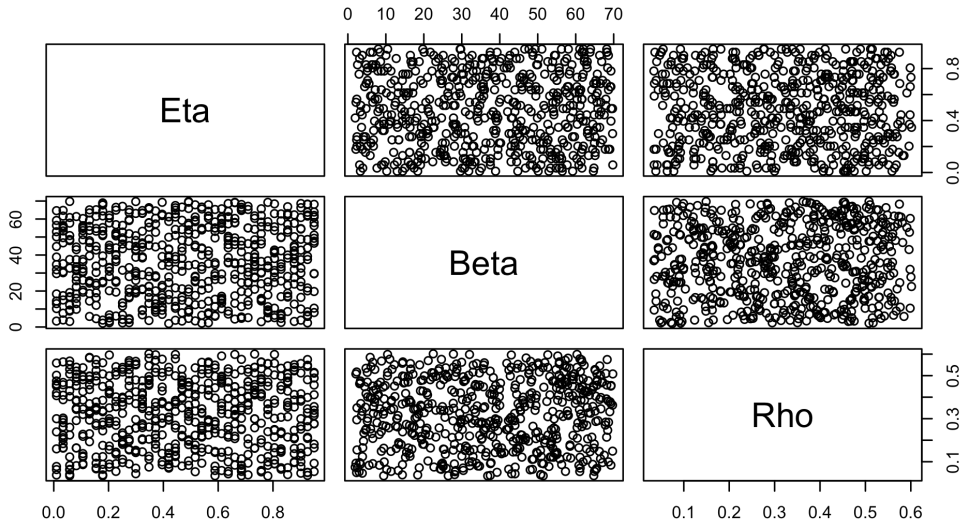
	Beta	Gamma	Rho	k	Eta	N
[1,]	1.824688	2	0.03405657	10	0.03628984	38000
[2,]	69.791919	2	0.60343428	10	0.99982180	38000

	loglik	loglik.se
[1,]	-122.7423	0.01462075
[2,]	-104.2847	0.56960880

Within this box, we'll choose some random starting points

```
freeze(seed=1196696958,  
  profile_design(  
    Eta=seq(0.01,0.95,length=40),  
    lower=box[1,c("Beta","Rho")],  
    upper=box[2,c("Beta","Rho")],  
    nprof=15, type="runif"  
  )) -> guesses
```

```
plot(guesses)
```




Now, we'll start one independent sequence of iterated filtering operations from each of these points.

```
foreach(guess=iter(guesses,"row"), .combine=rbind,  
  .options.future=list(seed=830007657)  
) %dofuture% {  
  mf1 |>  
    mif2(params=c(guess,fixed_params),  
      rw.sd=rw_sd(Beta=0.02,Rho=0.02)) |>  
    mif2(Nmif=100,cooling.fraction.50=0.3) -> mf  
  replicate(  
    10,  
    mf |> pfilter(Np=5000) |> logLik()) |>  
    logmeanexp(se=TRUE) -> ll  
  mf |> coef() |> bind_rows() |>  
    bind_cols(loglik=ll[1],loglik.se=ll[2])  
} -> results
```

Visualizing profile likelihood


```
read_csv("measles_params.csv") |> filter(loglik>max(loglik)-10) -> all  
pairs(~loglik+Beta+Eta+Rho,data=all,pch=16)
```



Focusing on the top of the surface

```
results |>
  filter(is.finite(loglik)) |>
  group_by(round(Eta,5)) |>
  filter(rank(-loglik)<3) |>
  ungroup() |>
  filter(loglik>max(loglik)-20) |>
  ggplot(aes(x=Eta,y=loglik))+
  geom_point() + xlab(expression(eta)) -> plot
```

plot




estimate η with 95% CI using these data

```
maxloglik <- max(results$loglik,na.rm=TRUE)
ci.cutoff <- maxloglik-0.5*qchisq(df=1,p=0.95)

results |>
  filter(is.finite(loglik)) |>
  group_by(round(Eta,5)) |>
  filter(rank(-loglik)<3) |>
  ungroup() |>
  ggplot(aes(x=Eta,y=loglik))+
  geom_point() + xlab(expression(eta)) +
  geom_smooth(method="loess",span=0.25)+
  geom_hline(color="red",yintercept=ci.cutoff)+
  lims(y=maxloglik-c(5,0)) -> plot
```

plot




```
results |>
  filter(is.finite(loglik)) |>
  filter(loglik>max(loglik)-0.5*qchisq(df=1,p=0.95)) |>
  summarize(min=min(Eta),max=max(Eta)) -> Eta_ci
```

Then we know η is in the 0.23–0.95% range (95% CI).

For the later parameters

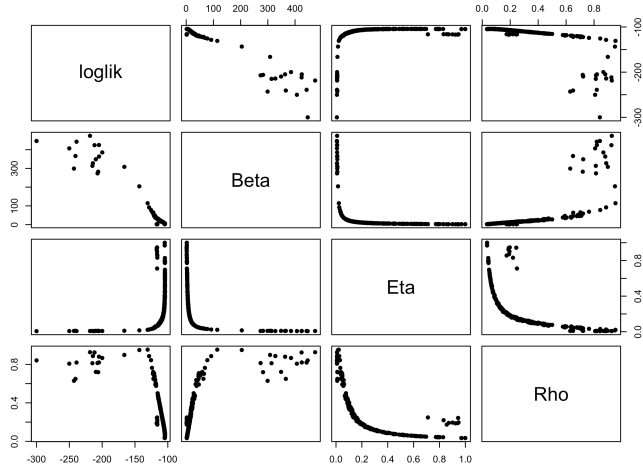
for next parameter, we can initialize the IF2 computations at points we have already established have high likelihoods.


```
read_csv("measles_params.csv") |>
  group_by(cut=round(Rho,2)) |>
  filter(rank(-loglik)<=10) |>
  ungroup() |>
  arrange(-loglik) |>
  select(-cut,-loglik,-loglik.se) -> guesses
```

then again, generate the profile likelihood

```
foreach(guess=iter(guesses,"row"), .combine=rbind,  
  .options.future=list(seed=2105684752)  
) %dofuture% {  
  mf1 |>  
    mif2(params=guess,  
      rw.sd=rw_sd(Beta=0.02,Eta=ivp(0.02))) |>  
    mif2(Nmif=100,cooling.fraction.50=0.3) |>  
    mif2() -> mf  
  replicate(  
    10,  
    mf |> pfilter(Np=5000) |> logLik()) |>  
    logmeanexp(se=TRUE) -> ll  
  mf |> coef() |> bind_rows() |>  
    bind_cols(loglik=ll[1],loglik.se=ll[2])  
} -> results
```

Visualizing profile likelihood





```
results |>  
  filter(loglik>max(loglik)-0.5*qchisq(df=1,p=0.95)) |>  
  summarize(min=min(Rho),max=max(Rho)) -> rho_ci
```

Then we know reporting efficiencies ρ is in the 3.4–16% range (95% CI).

The investigation continues....

Check predictions I

After all these analyses, we would like to visualize how exactly the model with the MLEs matches the data. We can do it by plotting the simulations with 95% the prediction interval.

```
read_csv("measles_params.csv") |>
  filter(loglik == max(loglik)) |>
  select(-loglik, -loglik.se) -> best.params

measSIR |>
  simulate(
    params=unlist(best.params),
    nsim=1000, format="data.frame", include.data=TRUE
  ) -> sims
```


Check predictions II

```
sims |>
  mutate(data=.id=="data") |>
  group_by(week,data) |>
  reframe(
    p=c(0.025,0.5,0.975),
    value=wquant(reports,probs=p),
    name=c("lo","med","up")
  ) |>
  select(-p) |> pivot_wider() |> ungroup() |>
  ggplot(aes(x=week,y=med,color=data,fill=data,ymin=lo,ymax=up))+
  geom_line()+ geom_ribbon(alpha=0.2,color=NA) +
  labs(y="reports")+
  theme_bw() + guides(color="none",fill="none")
```

Check predictions III

